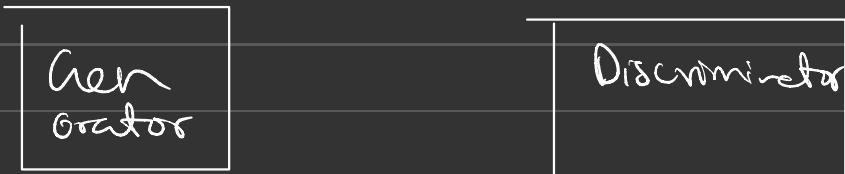



General Adversarial

Networks

Noise Class Factors
 ξ , $Y \rightarrow X$

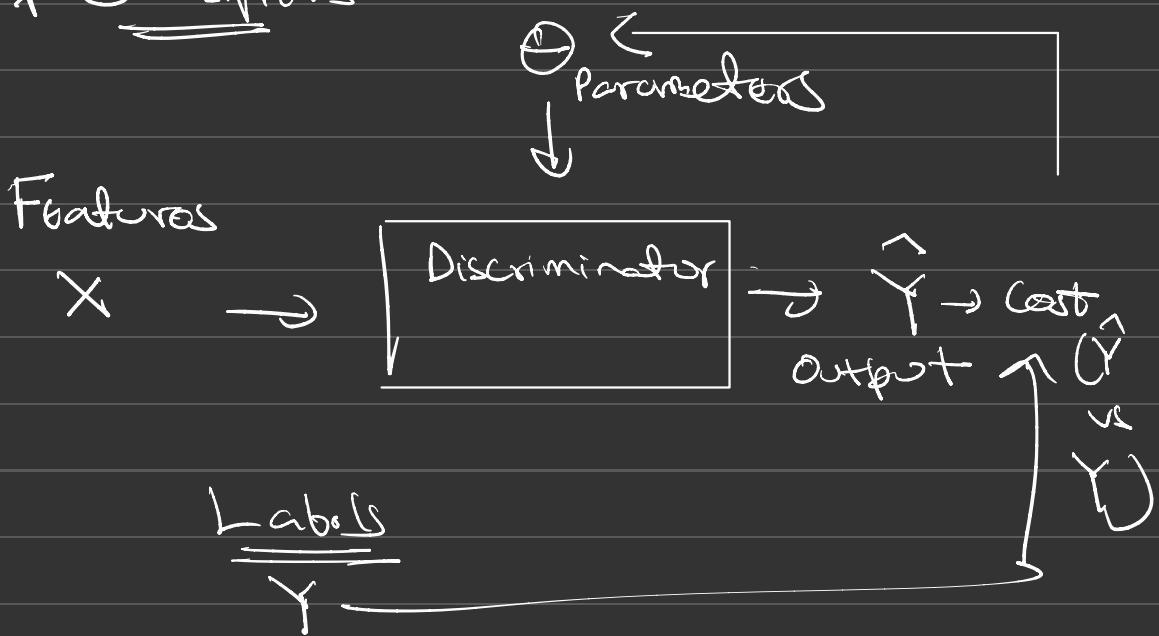
$$P(X|Y)$$



↳ Similar
to Decoder
in VAE

→ ~~Labels~~, we don't need the discriminator and we can use the generator.

~~Classifiers~~



Discriminator

$$P(Y|X)$$

Class Features

Generator

$$P(X|Y)$$

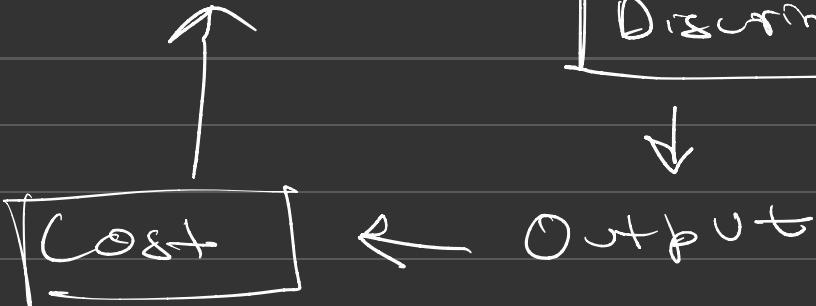
Features Class

Generator



θ
Parameters

Discriminator



★ Binary Cross Entropy

Average loss of
the whole batch

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \log h(x_i^\top \theta) + (1-y^{(i)}) \log (1-h(x_i^\top \theta)) \right]$$

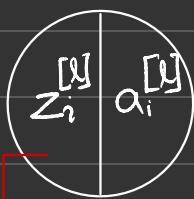
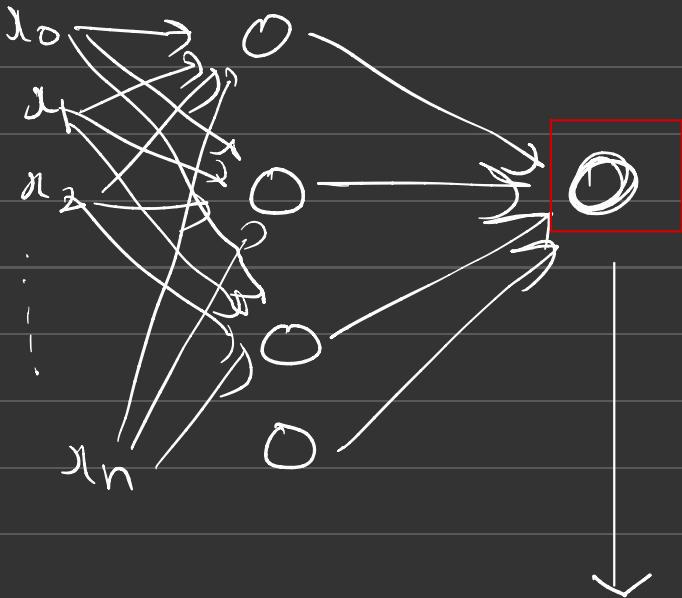
Prediction

Label

Features

Labels

→ We can't use a "superior" generator with a normal discriminator and vice versa



$$z_i^l = \sum_{i=0}^n w_i^{[l]} a_i^{[l-1]}$$

$$a_i^{[l]} = g^{[l]}(z_i^{[l]})$$

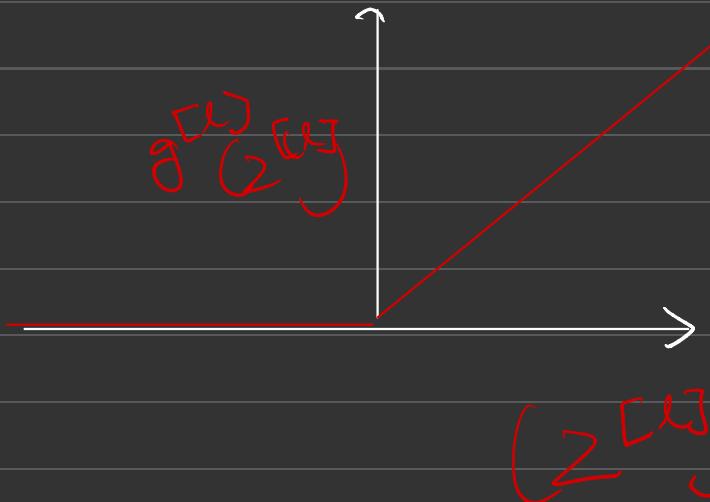
Diffs for backprop.

Non-l: non → for approx.
complex
function

activation
function

(Diff and Non-linear)

* ReLU



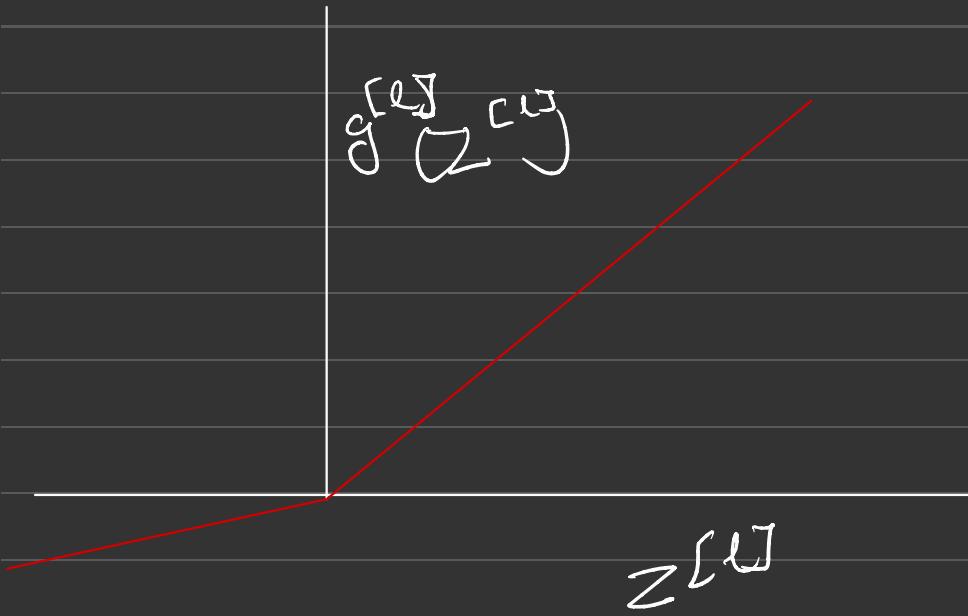
$$g^{[l]}(z^{[l]}) = \max(0, z^{[l]})$$

* Dying ReLU \rightarrow (due to the fact
per not accomodating
to the change)



Solution: Leaky ReLU

* Leaky ReLU

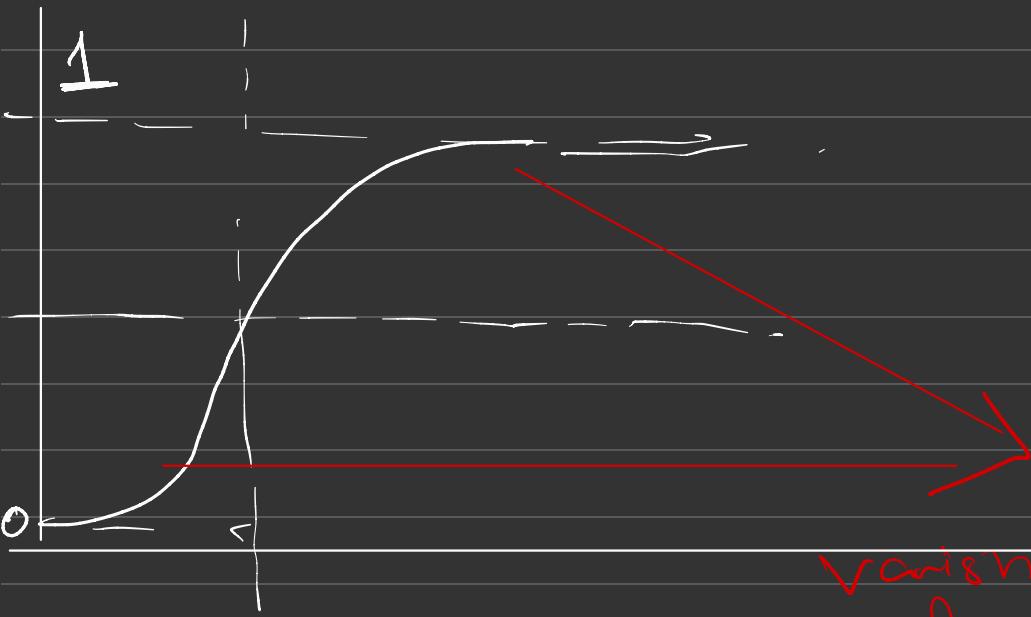


$$g^{[l]}(z^{[l]}) = \max(\alpha z^{[l]}, z^{[l]})$$

\downarrow
Solves
the dying ReLU
problem

Value often set to
0.1

* Sigmoid activation



vanishing
gradient
problem

★ Batch Normalisation

- to counter covariate shift
- To create similar distribution of features of factors that will result in a smoother cost function
- speeds up the learning process

★ Batch normalisation : training

$$z_i^{[l]} = \sum_{i=0} w_i^{[l]} a_i^{[l-1]}$$

$$\begin{array}{|c|c|}\hline z_i^{[l]} & a_i^{[l-1]} \\ \hline\end{array}$$

$$\hat{z}_i^{[l]} = \frac{z_i^{[l]} - \bar{m}_{z_i}^{[l]}}{\sqrt{o_{z_i}^2 + \epsilon}}$$

mean
 → to zero
 denom ≠ 0
 variance

$$y_i^{[l]} = \gamma \hat{z}_i^{[l]} + \phi$$

Scale ↓ Shift
 Factor Factor

* Batch Normalization : To fit

$$\hat{z}_i^{[l]} = \frac{z_i^{[l]} - E(z_i^{[l]})}{\sqrt{\text{Var}(z_i^{[l]}) + \epsilon}}$$

(now these
are fixed values) $\sqrt{\text{Var}(z_i^{[l]}) + \epsilon}$

★ Convolution

Stride → how many pixels you move in one time with the filter
(no. of steps basically)

Padding → extra frame of pixels used to pay attention to all the parts of the image

Pooling → To reduce the size of the images

★ Upsampling

↳ nearest neighbours



20	15
10	40

Upsampling
→

20	20	15	15
20	20	15	15
10	10	40	40
10	10	40	40

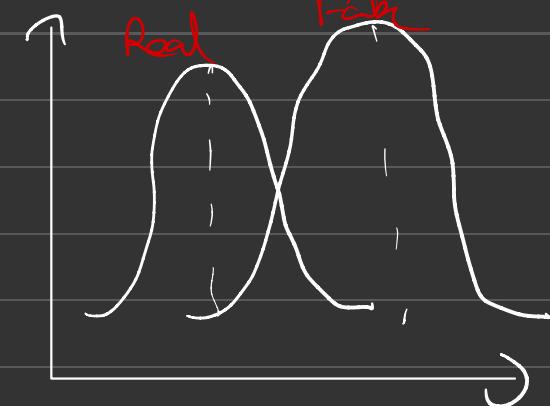
→ Pooling & upsampling do not have any learnable parameters.

↳ But transposed convolutions do!!

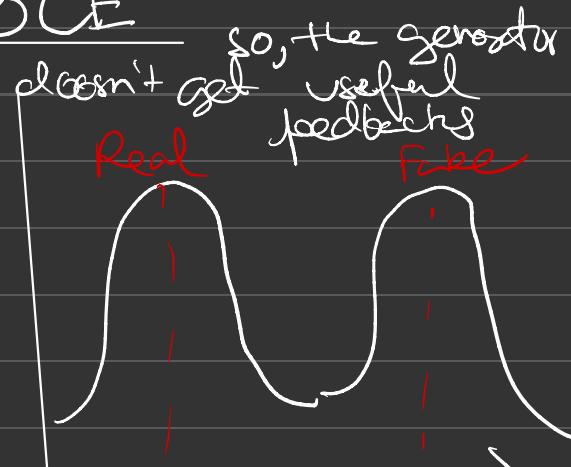
* Mode collapse

- Any peak on the density function is a mode
- Mode collapse happens when the generator gets stuck in one mode

* Problem with BCE

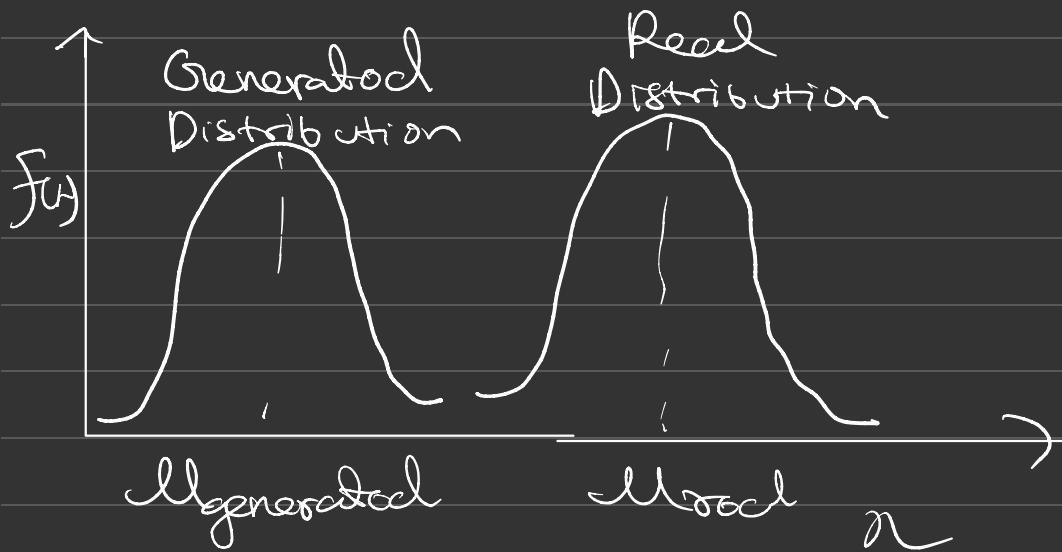


Discriminator
in the beginning



Discriminator
Tutor

★ Earth Mover's Distance



- It calculates the "Effort" to make the generated distribution equal to the real distribution
- Solves gradient vanishing problem

★ Wasserstein Loss

$$\min_g \max_c E(c_{\text{a}}) - E(c(g(\omega)))$$

BCE	W-Loss
Out put $0 < \alpha < 1$	Out put any number

→ also help with mode collapse

★ Condition of W-Loss

Critic should be 1-Lipschitz
continuous

The norm of the gradient
Should be at-most 1 for
every point

★ 1-Lipschitz cont.

$$\|\nabla f(w)\|_2 \leq 1$$

How to solve it?

→ Weight clipping forces the weights of the critic to a fixed interval

Gradient descent to update weights



Clip the critic's weight

(Limits the learning ability)
~~ability~~ of the critic

(bottom method)

★ Gradient Penalty

W-loss

Function



$$\left[\min_g \max_c E(c(\omega)) - E(c(g(\omega))) \right]$$

$$+ \boxed{\lambda \nabla g}$$



regularisation of
the critic's gradient

Final form of W-loss
→ the main Function

$$\min_g \max_c E(c(\omega)) - E(c(g(\omega)))$$

$$+ \lambda E\left(\|\nabla c(\hat{\omega})\|_2 - 1\right)^2$$

→ makes GAN less prone to mode collapse and vanishing gradient

→ tries to make the critic 1-L (ent. and drift)

* Conditional Generation

→ Training dataset needs to be labeled.



Class (one)

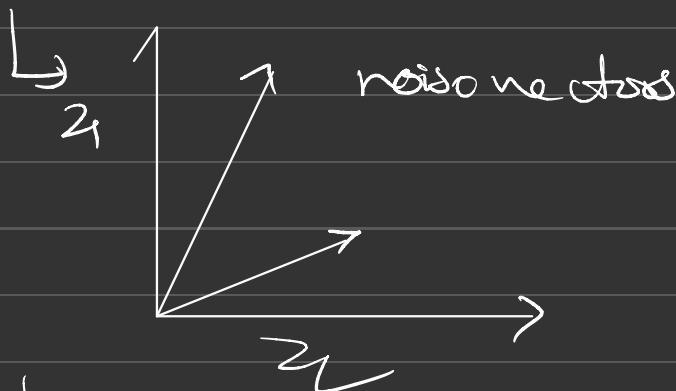
$$\text{hot-} \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Vector

* Controllable GAN

- Control by changing the noise vector
- does not need a labeled training dataset

→ Z-space

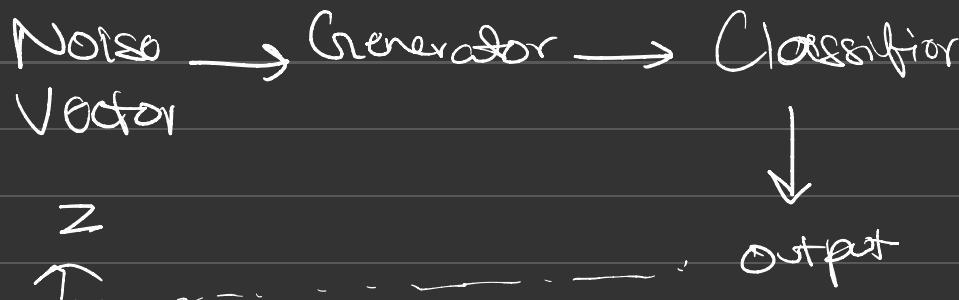


↳ problem:-

↳ Z-space Entanglement

* Classifier Gradients

Protrained



Direction to make
the images look like
they have sunglasses